

```
In [25]: import numpy as np
import pandas as pd
```

```
In [26]: data = pd.read_csv("HousingData.csv")
data.head()
```

```
Out[26]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	

```
In [27]: data.tail()
```

```
Out[27]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	21.0	396.90	

```
In [28]: print("The shape of the data is: ")
data.shape
```

The shape of the data is:

```
Out[28]: (506, 14)
```

```
In [29]: data.isnull().sum()
```

```
Out[29]: CRIM      20
ZN          20
INDUS       20
CHAS        20
NOX         0
RM          0
AGE         20
DIS         0
RAD         0
TAX         0
PTRATIO     0
B           0
LSTAT       20
MEDV        0
dtype: int64
```

```
In [30]: X = data.iloc[:,0:13]
        y = data.iloc[:, -1]
```

```
In [31]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, ra
```

```
In [32]: print(X_train.shape)
        print(X_test.shape)
        print(y_train.shape)
        print(y_test.shape)
```

```
(404, 13)
(102, 13)
(404,)
(102,)
```

```
In [33]: from sklearn.linear_model import LinearRegression
```

```
In [37]: from sklearn.impute import SimpleImputer
        from sklearn.preprocessing import StandardScaler
        from sklearn.linear_model import LinearRegression
        from sklearn.pipeline import make_pipeline

        # Create a pipeline that includes imputation, scaling, and regression
        model = make_pipeline(
            SimpleImputer(strategy='mean'), # Replace 'mean' with 'median', 'most_
            StandardScaler(with_mean=False),
            LinearRegression()
        )

        model.fit(X_train, y_train)
```

```
Out[37]: Pipeline(steps=[('simpleimputer', SimpleImputer()),
                          ('standardscaler', StandardScaler(with_mean=False)),
                          ('linearregression', LinearRegression())])
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [40]: model.score(X_test, y_test)
```

```
Out[40]: 0.6590604241860214
```

```
In [ ]:
```

```
In [ ]:
```

