

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [3]: df= pd.read_csv('Social_Network_Ads.csv')
```

```
In [4]: df.shape
```

```
Out[4]: (400, 5)
```

```
In [5]: df.head(15)
```

```
Out[5]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0
10	15570769	Female	26	80000	0
11	15606274	Female	26	52000	0
12	15746139	Male	20	86000	0
13	15704987	Male	32	18000	0
14	15628972	Male	18	82000	0

```
In [7]: df.drop(['User ID'],axis=1,inplace=True)
```

```
In [8]: df.head()
```

```
Out[8]:
```

	Gender	Age	EstimatedSalary	Purchased
0	Male	19	19000	0
1	Male	35	20000	0
2	Female	26	43000	0
3	Female	27	57000	0
4	Male	19	76000	0

```
In [9]: df.Purchased.value_counts()
```

```
Out[9]: Purchased
0      257
1      143
Name: count, dtype: int64
```

```
In [10]: df.Gender.value_counts()
```

```
Out[10]: Gender
Female    204
Male      196
Name: count, dtype: int64
```

```
In [11]: df.dtypes
```

```
Out[11]: Gender          object
Age              int64
EstimatedSalary  int64
Purchased        int64
dtype: object
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: Gender          0
Age              0
EstimatedSalary  0
Purchased        0
dtype: int64
```

```
In [13]: df.describe()
```


```
Out[13]:
```

	Age	EstimatedSalary	Purchased
<b>count</b>	400.000000	400.000000	400.000000
<b>mean</b>	37.655000	69742.500000	0.357500
<b>std</b>	10.482877	34096.960282	0.479864
<b>min</b>	18.000000	15000.000000	0.000000
<b>25%</b>	29.750000	43000.000000	0.000000
<b>50%</b>	37.000000	70000.000000	0.000000
<b>75%</b>	46.000000	88000.000000	1.000000
<b>max</b>	60.000000	150000.000000	1.000000

```
In [14]: g = sns.catplot(x = "Gender",y = "Purchased",data = df,kind = "bar",height
g.set_y_labels("Purchased Probability")
plt.show
```

C:\Users\Aniket\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)

```
Out[14]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [15]: X=df.drop(['Gender','Purchased'],axis=1)
Y= df['Purchased']
X.head()
```

```
Out[15]:
```

	Age	EstimatedSalary
0	19	19000
1	35	20000
2	26	43000
3	27	57000
4	19	76000

```
In [18]: from sklearn.model_selection import train_test_split
# Shuffle and split the data into training and testing subsets
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, ra

# Success
print("Training and testing split was successful.")
```

Training and testing split was successful.

```
In [19]: from sklearn.linear_model import LogisticRegression
basemodel= LogisticRegression()
basemodel.fit(X_train,y_train)
print("Training accuracy:", basemodel.score(X_train,y_train)*100)
```

Training accuracy: 64.0625

```
In [20]: y_predict= basemodel.predict(X_test)
print("Testing accuracy:", basemodel.score(X_test,y_test)*100)
```

Testing accuracy: 65.0

```
In [22]: from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
```

```
In [23]: X=df[['Age','EstimatedSalary']]
X_scaled= scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, Y, test_size=
print("Training and testing split was successful.")
```

Training and testing split was successful.

```
In [24]: model= LogisticRegression()
model.fit(X_train,y_train)
y_predict= model.predict(X_test)
print("Training accuracy:", model.score(X_train,y_train)*100)
print("Testing accuracy:", model.score(X_test,y_test)*100)
```

Training accuracy: 80.9375

Testing accuracy: 87.5

```
In [25]: from sklearn.metrics import accuracy_score
Acc=accuracy_score(y_test,y_predict)
print(Acc)
```

0.875

```
In [26]: from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test,y_predict)
print(cm)
```

```
[[51  1]
 [ 9 19]]
```

```
In [27]: from sklearn.metrics import precision_recall_fscore_support
prf= precision_recall_fscore_support(y_test,y_predict)
print('precision:',prf[0])
print('Recall:',prf[1])
print('fscore:',prf[2])
print('support:',prf[3])
```

```
precision: [0.85 0.95]
Recall: [0.98076923 0.67857143]
fscore: [0.91071429 0.79166667]
support: [52 28]
```

```
In [28]: from sklearn.metrics import classification_report
cr= classification_report(y_test,y_predict)
print(cr)
```

	precision	recall	f1-score	support
0	0.85	0.98	0.91	52
1	0.95	0.68	0.79	28
accuracy			0.88	80
macro avg	0.90	0.83	0.85	80
weighted avg	0.89	0.88	0.87	80

```
In [ ]:
```